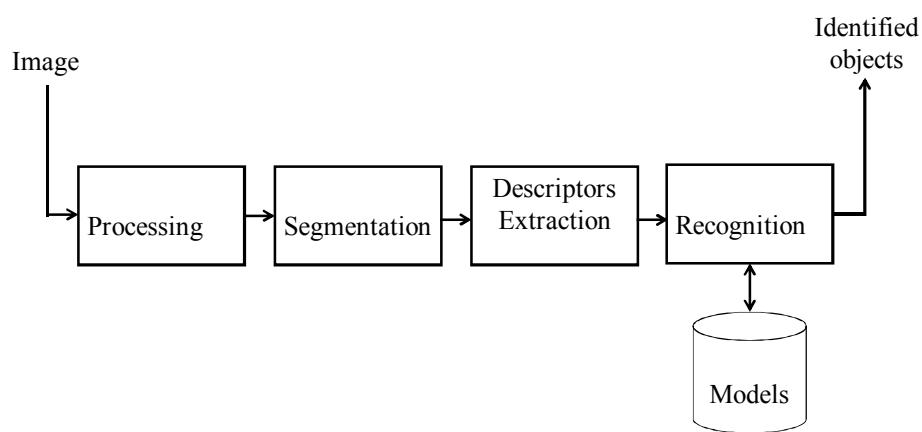


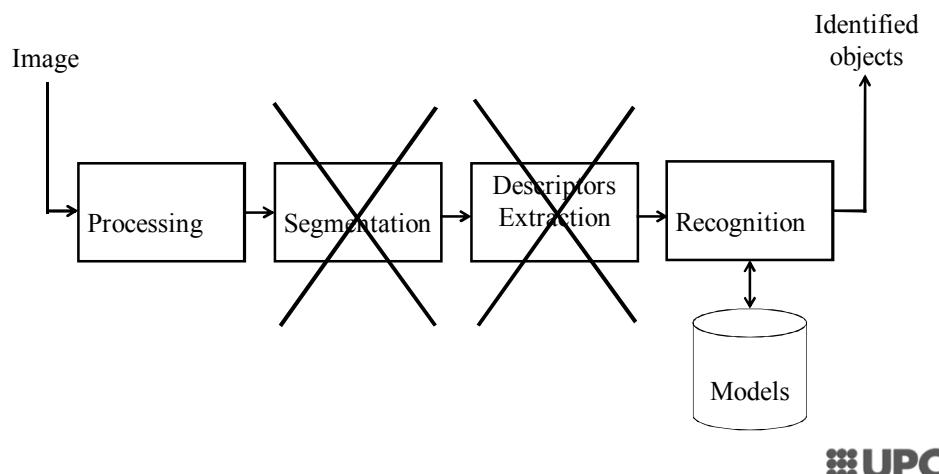
Local Features



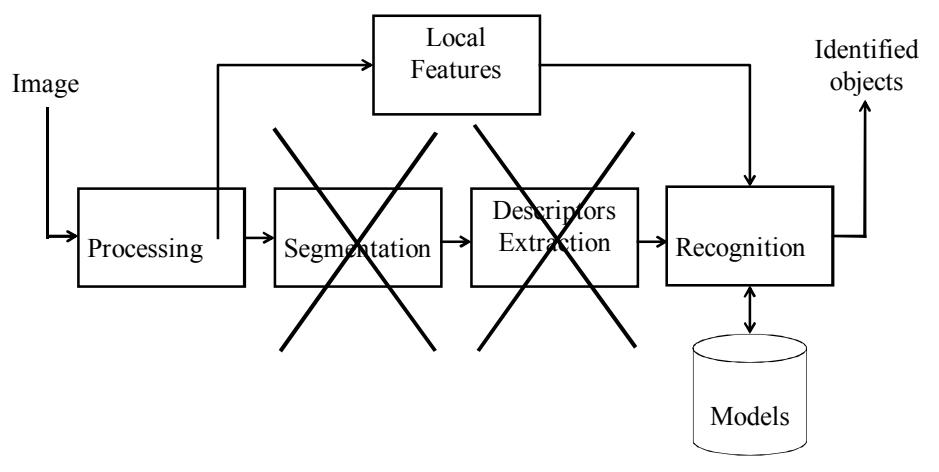
A computer vision system



A computer vision system



A computer vision system



Local Features

- Histograms
- Hough transform
- Corners
- Scale Invariant Feature transform (SIFT)
- Local Binary Patterns (LBP)
- Haar Features (face detection)



Local Features

- Global features are sensitive to noise and occlusion.
- An alternative is to use local features:
 - If some features are occluded, the object might still be recognized by identifying some visible features



Corners are Simple Features

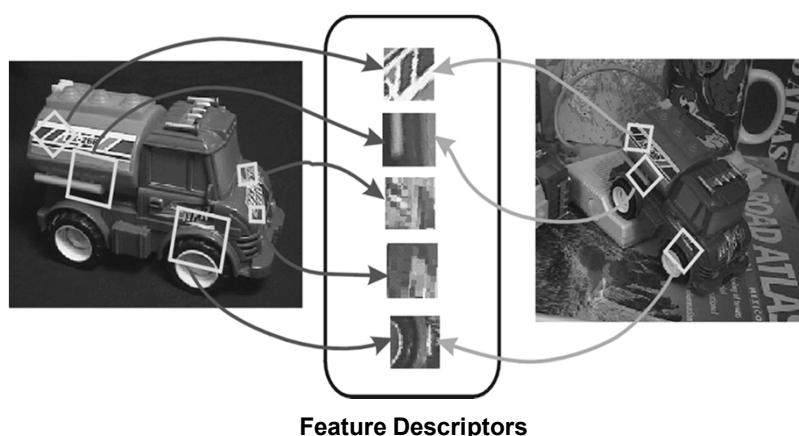
↳ Can you see common feature point between the two images?



Invariant local features

Find features that are invariant to transformations

- geometric invariance: translation, rotation, scale
- photometric invariance: brightness, exposure, ...



Advantages of local features

Locality

- features are local, so robust to occlusion and clutter

Distinctiveness:

- can differentiate a large database of objects

Quantity

- hundreds or thousands in a single image

Efficiency

- real-time performance achievable

Generality

- exploit different types of features in different situations

Properties of ‘good’ features

↳ Repeatability

- Should be detectable at the same locations in different images despite changes in viewpoint and illumination

↳ Saliency (descriptiveness)

- Same points in different images should have similar features
- And vice-versa

↳ Compactness (affects speed of matching)

- Fewer features
- Smaller features

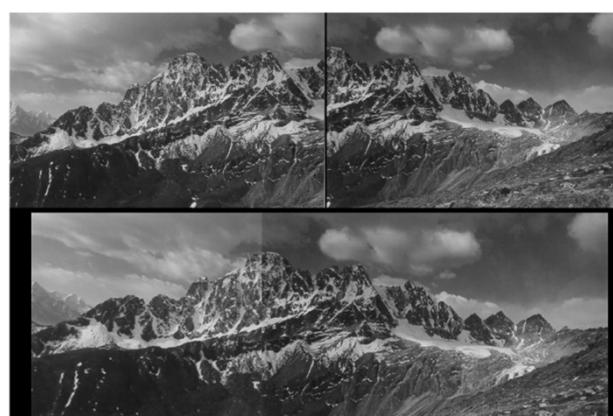
Applications of feature matching

- Image alignment and stitching
- Image retrieval
- Motion tracking
- Robot navigation
- Face detection
- Object recognition
- Gesture recognition
- Human action understanding
- Biometric identification



Applications of feature matching

Image Stitching (to generate panorama)



Applications of feature matching



■ UPC

Histograms

■ UPC

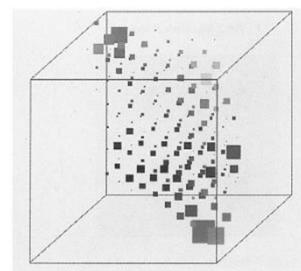
Colour histograms

- ↳ Colour stays constant under geometric transformations
- ↳ Colour is a local feature
 - It is defined for each pixel
 - It is robust to partial occlusion
- ↳ Idea:
 - can use object colours directly for recognition, or
 - better – use statistics of object colours
- ↳ Colour histogram is a type of appearance features



Colour histograms

- ↳ Colour histograms are colour statistics
 - Here, RGB as an example
 - Given: tristimulus R, G, B for each pixel
 - Compute a 3D histogram
 - $h(R, G, B) = \#(\text{pixels with colour } (R, G, B))$



Colour histograms

Colour Normalization

- ↳ One component of the 3D colour space is intensity
 - If a colour vector is multiplied by a scalar, the intensity changes but not the colour itself.
 - This means colours can be normalized by the intensity.
 - Note: intensity is given by $I = (R + G + B)/3$
 - Chromatic representation:

$$r = \frac{R}{R + G + B} \quad g = \frac{G}{R + G + B} \quad b = \frac{B}{R + G + B}$$

Since $r + g + b = 1$, only 2 parameters are needed to represent colour (knowing r and g , we can deduce $b = 1 - r - g$).

⇒ Can compute colour histogram using r , g , and b instead.

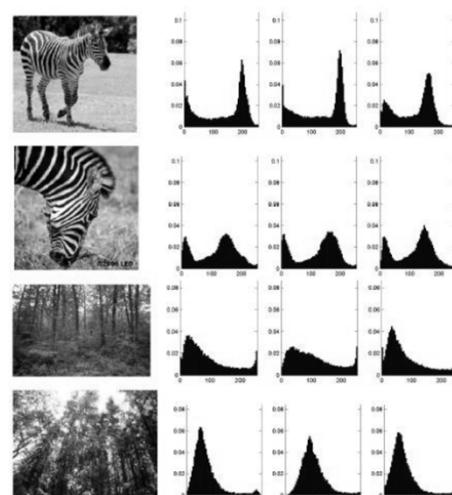


Colour histograms

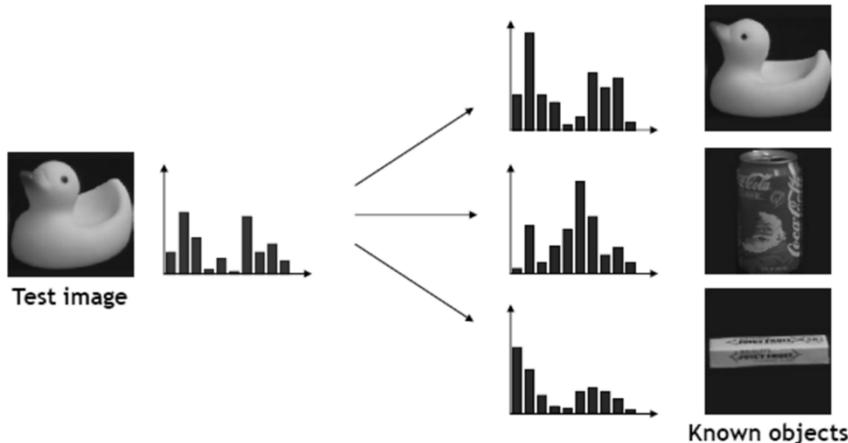
Object Recognition based on Colour Histograms

Colour histograms

- are discrete approximation of the colour distribution of an image.
- contain no spatial information ⇒ invariant to translation, scale, and rotation



Colour histograms



Object recognition using histograms

Simple algorithm

1. Build a set of histograms $H = \{h_i\}$ for each known object.
 - More exactly, for each view of each object.
2. Build a histogram h_t for the test image.
3. Compare h_t with each $h_i \in H$ using a suitable histogram comparison measure.
4. Select the object with the best matching score;
or reject the test image if no object is similar enough.

This is known as the “nearest-neighbour” strategy.



Colour histograms

Histogram Comparison with Multiple Training Views

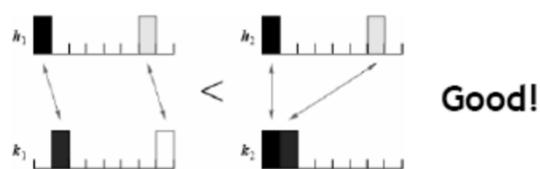
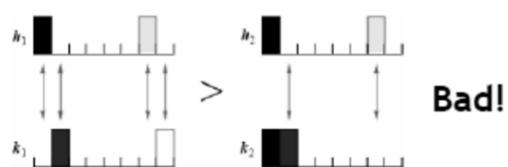
⇒ Need a good comparison measure for colour histograms !



Histogram comparison

What is a Good Comparison Measure?

- ↳ How to define matching cost?

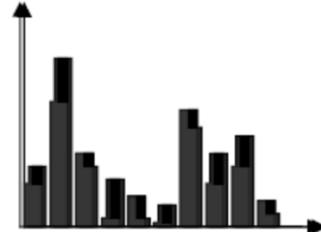


Histogram comparison

Euclidean distance (L_2 norm)

$$d(\mathbf{q}, \mathbf{v}) = \sum_i (q_i - v_i)^2$$

- ↳ Motivation of the Euclidean distance:
 - Focuses on the differences between the histograms.
 - Interpretation: distance in the feature space.
 - Range: $[0, \infty)$.
 - All cells are weighted equally.
 - Not very robust to outliers !

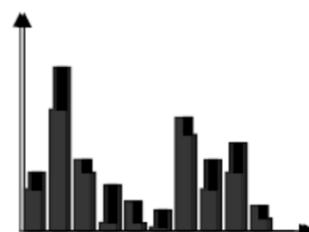


Histogram comparison

Chi-Square distance:

$$d(\mathbf{q}, \mathbf{v}) = \sum_i \frac{(q_i - v_i)^2}{q_i + v_i}$$

- ↳ Motivation of the χ^2 distance:
 - Statistical background
 - Test if two distributions are different.
 - Possible to compute a significance score.
 - Range: $[0, \infty)$.
 - Cells are not weighted equally !
 - More robust to outliers than the Euclidean distance, if the histograms contain enough observations...



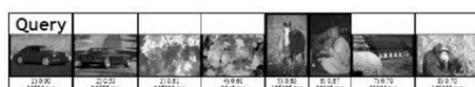
Histogram comparison

- ↳ Which measure is the best?
 - It depends on the application
 - Euclidean distance is often not robust enough.
 - Generally, χ^2 distance gives good performance for histograms
 - KL/Jeffreys divergence works well sometimes, but is expensive
 - EMD is the most powerful, but also very expensive.

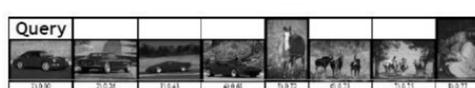


Colour histograms for image retrieval

- ↳ The image retrieval problem concerns the retrieval of those images in a database that best match a query image.



L2 distance



Jeffrey divergence

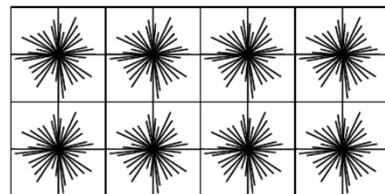
 χ^2 distance

Earth Movers Distance



Histograms of oriented gradients (HOGs)

- Objective: object recognition
- Basic idea
 - Local shape information often well described by the distribution of intensity gradients or edge directions even without precise information about the location of the edges themselves.



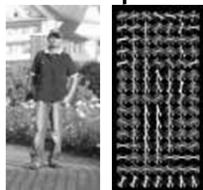
HOG algorithm

- Divide image into small sub-images: “cells”
- Accumulate a histogram of edge orientations within that cell
- The combined histogram entries are used as the feature vector describing the object
- To provide better illumination invariance (lighting, shadows,etc.) normalize the cells across larger regions incorporating multiple cells: “blocks”

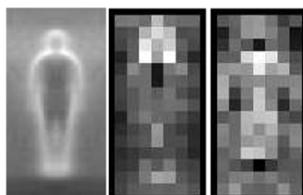


Person detection, Dalal-Triggs 2005

1. Represent each example with a single, fixed HoG template



2. Learn a single [linear] SVM as a detector



Code available: <http://pascal.inrialpes.fr/soft/olt/>

Positive and negative examples



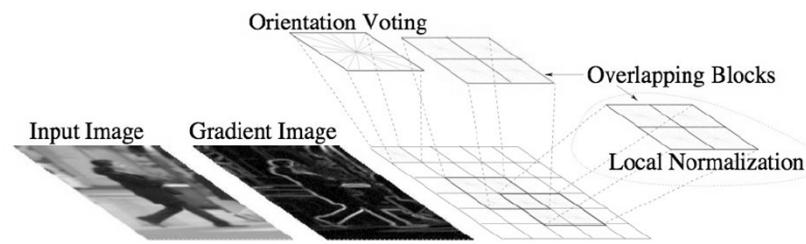
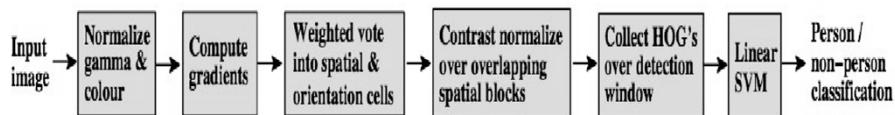
+ thousands more...



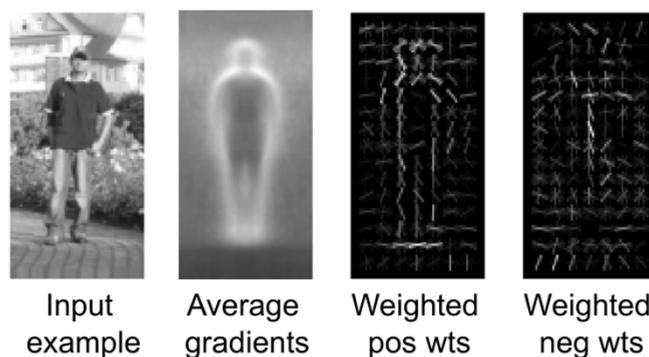
+ millions more...

HOGs. Aplicacions

A worked example: Dalal and Triggs CVPR'05



HOGs. Descriptor cues

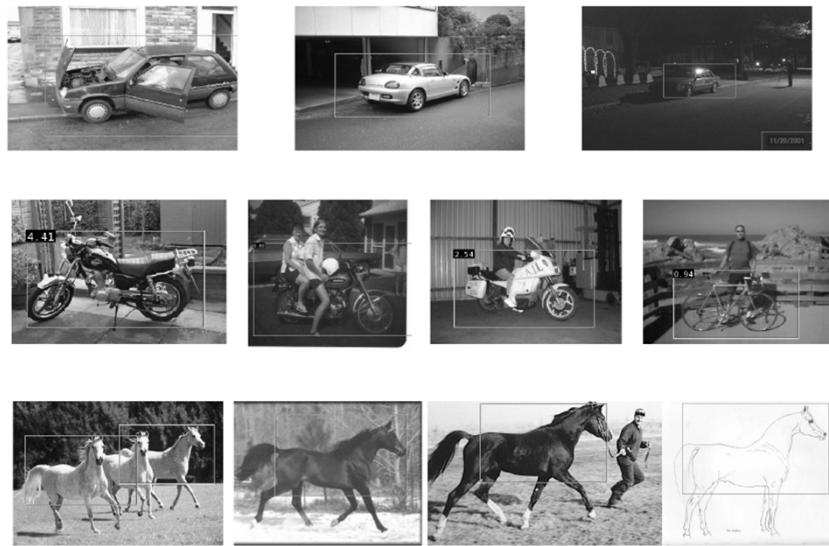


Most important cues are head, shoulder, leg silhouettes

Vertical gradients inside a person are counted as negative



Applications to Other Classes

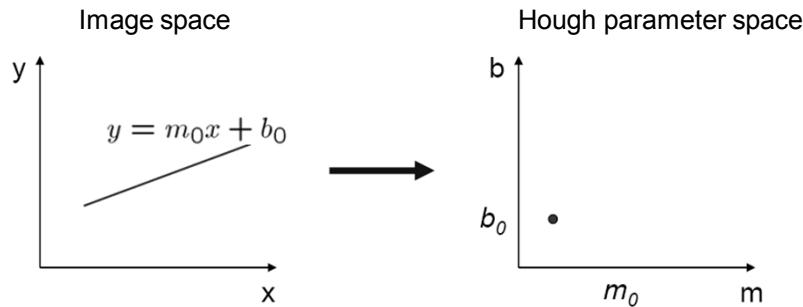


M. Everingham et al. *The 2005 PASCAL Visual Object Classes Challenge*. Proceedings of the PASCAL Challenge Workshop, 2006. 25

Transformada de Hough

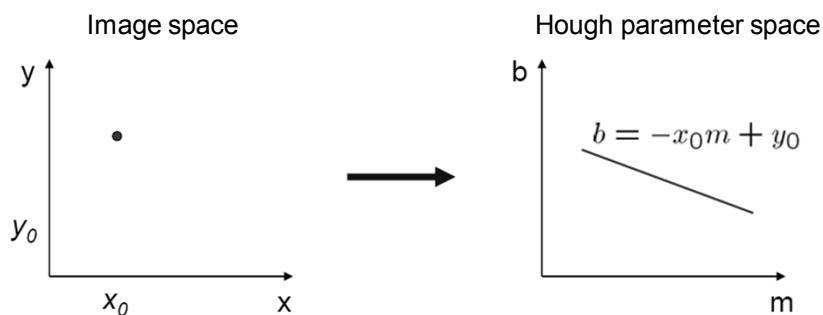
Transformada de Hough

- Dissenyada originalment per a la detecció de rectes
- Molt robusta a soroll i imperfeccions
- Cada punt de la recta original es transforma en una recta en el pla transformat



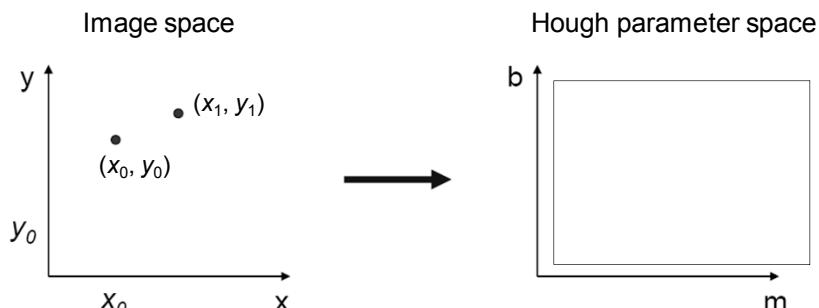
Parameter space representation

- What does a point (x_0, y_0) in the image space map to in the Hough space?
 - Answer: the solutions of $b = -x_0m + y_0$
 - This is a line in Hough space



Parameter space representation

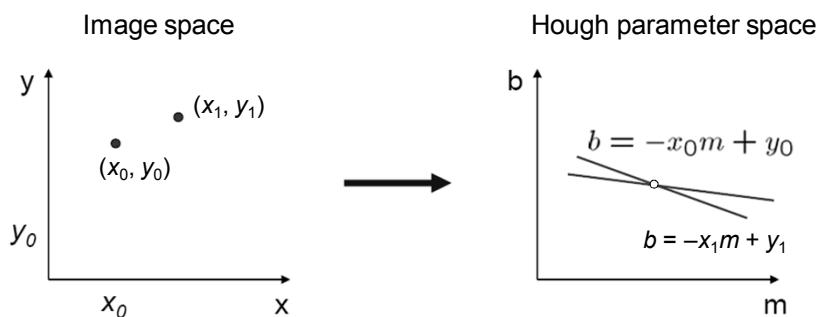
- Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?



Source: S. Seitz

Parameter space representation

- Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?
 - It is the intersection of the lines $b = -x_0m + y_0$ and $b = -x_1m + y_1$



Source: S. Seitz

Transformada de Hough

Algorisme:

- 1. Crear la matriu de l'espai m-b i dimensionar-la.
- 2. Cada posició de la matriu és un acumulador. Posar-los tots a 0.
- 3. Per a cada píxel (x,y) de la imatge, incrementar les posicions de la matriu m-b que satisfan l'equació.
- 4. Els màxims de la matriu m-b es corresponen a la presència de rectes en la imatge.

- Si hi ha diverses rectes a la imatge, tindrem diversos màxims a la matriu.
- Es difícil determinar si un màxim és determinant o no.



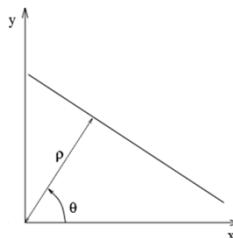
Transformada de Hough

-Problema: afitar m i b.

-Solució: S'usa la forma polar de la recta:

$$r = x \cos \alpha + y \sin \alpha$$

- ara els punts (x,y) queden mapejats com funcions sinusoïdals en l'espai (r,α) .
- r representa la llargada i α l'orientació del vector normal desde l'origen de la imatge a la línia.



Transformada de Hough

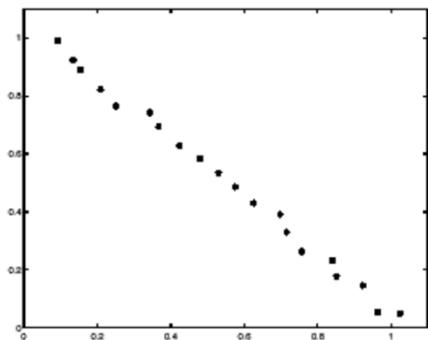
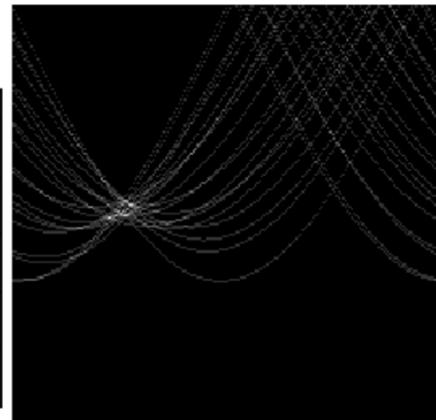


Image space

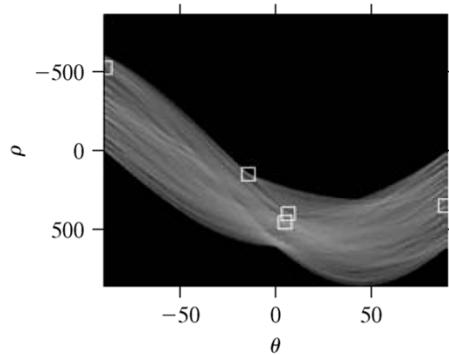


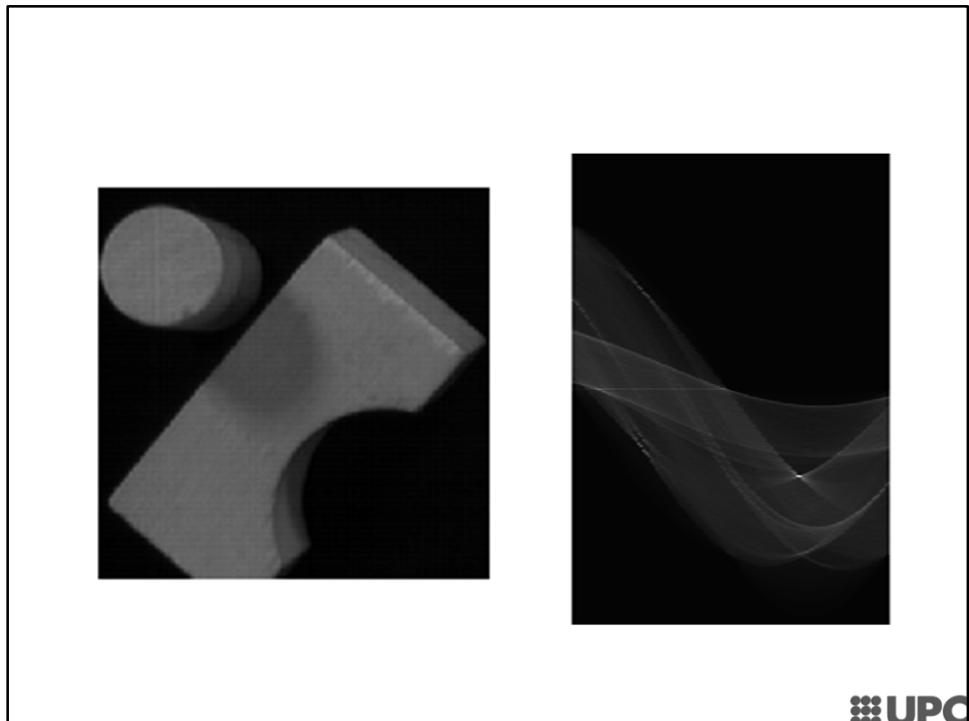
Votes

(Horizontal axis is θ , vertical is ρ)

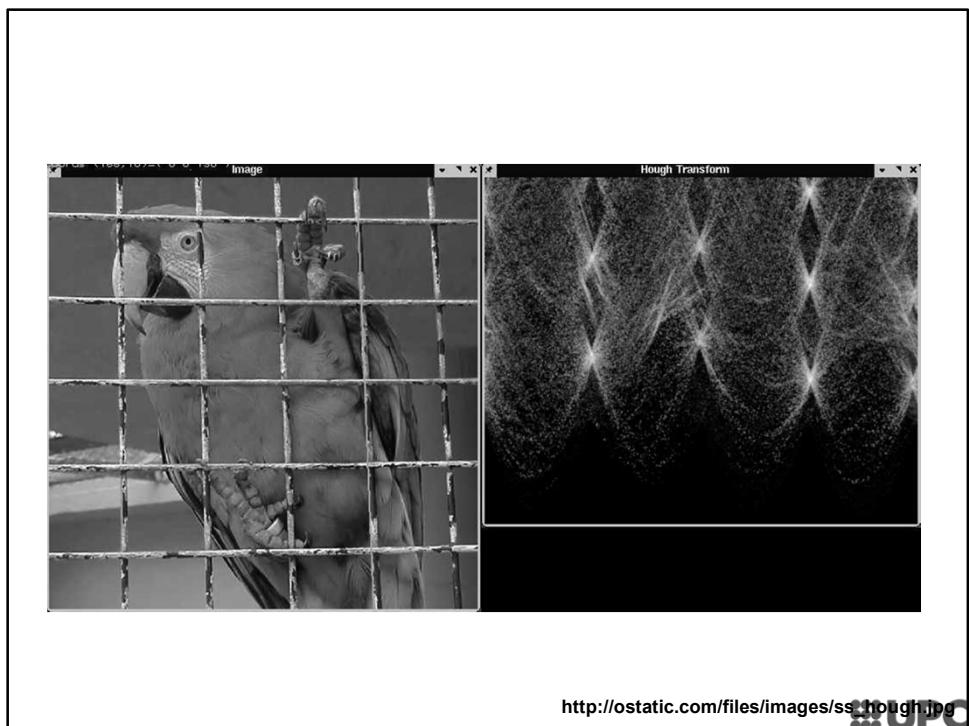


Transformada de Hough





■ UPC



http://ostatic.com/files/images/ss_hough.jpg

■ UPC

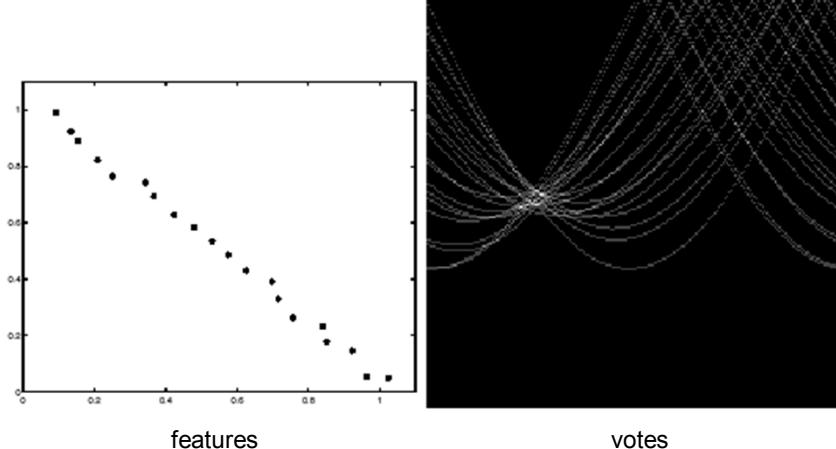
Transformada de Hough

Mechanics of the Hough transform

- Difficulties
 - how big should the cells be? (too big, and we merge quite different lines; too small, and noise causes lines to be missed)
 - small cell sizes => more computation!
 - large cell sizes => less accuracy!
- How many lines?
 - Count the peaks in the Hough array
 - Treat adjacent peaks as a single peak, i.e. must perform non-maximum suppression



Effect of noise



Peak gets fuzzy and hard to locate

Dealing with noise

- Choose a good grid / discretization
 - Too coarse: large votes obtained when too many different lines correspond to a single bucket
 - Too fine: miss lines because some points that are not exactly collinear cast votes for different buckets
- Increment neighboring bins (smoothing in accumulator array)
- Try to get rid of irrelevant features
 - Take only edge points with significant gradient magnitude

Hough Transform Algorithm

Input is an edge image ($E(i,j)=1$ for edgels)

1. Discretize θ and ρ in increments of d_θ and d_ρ .
Let $A(R,T)$ be an array of integer accumulators, initialized to 0.
2. For each pixel $E(i,j)=1$ and $h=1,2,\dots,T$ do
 1. $\rho = i \cos(h * d_\theta) + j \sin(h * d_\theta)$
 2. Find closest integer k corresponding to ρ
 3. Increment counter $A(h,k)$ by one
 3. Find local maxima in $A(R,T)$

Hough Transform Speed Up

- If we know the orientation of the edge - usually available from the edge detection step
 - We fix theta in the parameter space and increment **only one** counter!
 - We can allow for orientation uncertainty by incrementing a *few* counters around the "nominal" counter.



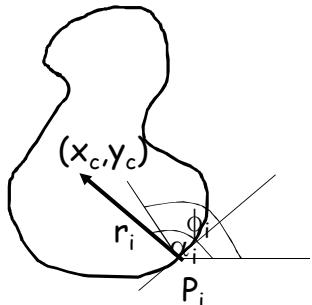
Hough Transform for Curves

- The H.T. can be generalized to detect any curve that can be expressed in parametric form:
 - $y = f(x, a_1, a_2, \dots, a_p)$
 - a_1, a_2, \dots, a_p are the parameters
 - The parameter space is p -dimensional
 - The accumulating array is **LARGE!**



Generalizing the H.T.

The H.T. can be used even if the curve has not a simple analytic form!



1. Pick a reference point (x_c, y_c)
2. For $i = 1, \dots, n$:
 1. Draw segment to P_i on the boundary.
 2. Measure its length r_i , and its orientation α_i .
 3. Write the coordinates of (x_c, y_c) as a function of r_i and α_i
 4. Record the gradient orientation ϕ_i at P_i .
3. Build a table with the data, indexed by ϕ_i .

$$x_c = x_i + r_i \cos(\alpha_i)$$

$$y_c = y_i + r_i \sin(\alpha_i)$$

PENNSTATE

Transformada de Hough

Finding Circles by Hough Transform

- ↳ Equation of Circle:

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- ↳ If radius is known: (2D Hough Space)

- Accumulator Array $A(a, b)$

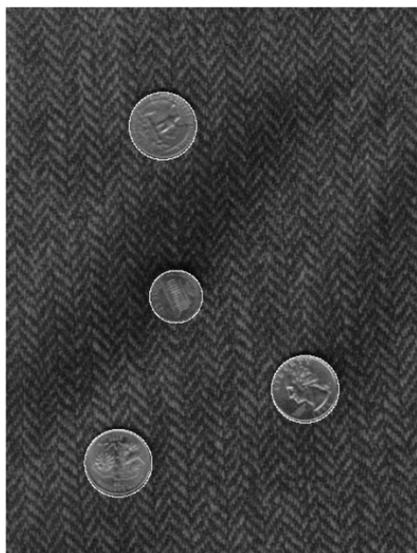
- ↳ If radius is not known: 3D Hough Space!

- Use Accumulator array $A(a, b, r)$

UPC

Transformada de Hough

Finding Coins



Note that because the quarters and penny are different sizes, a different Hough transform (with separate accumulators) was used for each circle size.